

Opportunity Radar

Where the build actually stands, what's confirmed working, and what phase 4 needs to know before starting.

Updated 21 September 2026

115

automated tests
passing

10

features added this
week

0.1.24

current build, live on
Apify

4/5

phases substantially
done

The honest headline

What's actually confirmed working

Scraping, eligibility matching, and trust/scam scoring have all been run live against real scholarship listings and produce correct, sensible results. This is the proven core.

What's built but not yet live-tested

10 additional features were built and unit tested today (115 tests passing, zero live API calls spent testing them, on purpose, to protect quota). None have been confirmed end to end against a real run yet, that's the next thing once today's Gemini quota resets.

Phase status

DONE

Phase 0–1 — Project setup and scraping. Gets past PhDportal's bot detection via Apify's Website Content Crawler, pulls real listings with full eligibility text from each listing's own detail page.

DONE

Phase 2 — Eligibility matching. Hard facts (deadline, nationality, education level) checked with plain code; an LLM only interprets genuinely ambiguous language.

Confirmed against a real run: a Nigerian profile was correctly rejected from an Austrian-only scholarship with no LLM call needed.

DONE

Phase 3 — Trust/scam scoring. Weighted red-flag detection plus real web-search evidence (via DuckDuckGo, since Google Custom Search and Bing's search APIs are both closed to new accounts). An LLM filters out false-positive keyword matches before they count as evidence.

PARTIAL

Phase 4 — Second source (Opportunity Desk) is coded and unit tested, but off by default and untested live. **Frontend hasn't started yet, this is where the team picks up.**

PARTIAL

Phase 5 — Eval sets done for both matching (7/7) and trust scoring (14/14), all passing. Demo script written. Live rehearsal hasn't happened yet.

The 10 features added this week

All isolated by design: if any one fails, the core matching and trust scoring keep working exactly as before. Built with mocked data to avoid spending real API quota, real confirmation still pending.

Plain-English digest

A one-line summary of results, e.g. "9 eligible, 10 need a closer look." ON

Deadline urgency

Flags listings as Closing soon / Upcoming / Plenty of time. ON

Trust confidence

Separate from risk level, how much evidence actually backed the verdict. ON

Action steps

Concrete next steps for Partial matches, not just "unclear." ON

Alumni signal

Checks for public mentions of real past recipients. OFF BY DEFAULT

CV-as-evidence

Optional CV text checked against ambiguous criteria instead of left unconfirmed. ON

Cross-listing fraud detection

Flags the same scam phrase repeating across multiple listings in one run. ON

Growing scam-pattern library

Learns new scam phrasing between runs, stored on Apify. ON

Compare-two-listings mode

Second source support

A separate input mode for a head-to-head between two opportunities. **ON**

Opportunity Desk scraping, code ready.

OFF BY DEFAULT

What phase 4 needs to know

The full input/output contract lives in `FRONTEND-SCHEMA.md` in the repo, with real examples pulled from actual runs. Start there, not from this document.

- The schema is stable, it won't change in breaking ways from here.
- A full run currently takes several minutes, design the loading state around that.
- `educationLevel` now includes "High school" alongside Bachelors/Masters/PhD.
- Compare-mode is a separate input shape, two already-scored listings in, a verdict out, instant, no scraping involved.

Known, honest limitations

Worth knowing rather than discovering during a demo:

- Gemini's free tier caps at 5 requests/minute and 20 requests/day per key, this shapes how much testing can happen in a day.
- DuckDuckGo search evidence does keyword matching; an LLM filters obviously irrelevant results, but it isn't a perfect relevance check.
- Opportunity Desk is Cloudflare-protected and has never been confirmed working against the live pipeline.

Opportunity Radar — built for She Code Africa x Apify BuildHER Hackathon 2026. Repo: github.com/LovableLynx/opportunity-radar (private).